

# RDMA Programming

- #include <rdma/rdma\_cma.h>
- #include <infiniband/verbs.h>

# RDMA Workflow

- Connection Setup
- Memory Registration
- Posting Work Requests (WRs)
- Completion Handling

# Connection Setup

- Connection Manager (CM): Manages connection setup and teardown

Open a channel to report asynchronous connection events -  
`rdma_create_event_channel()`

Create connection identifier (like socket) -  
`rdma_create_id(channel, &cm_id, context, RDMA_PS_TCP)`

- Client Resolve address and route: Map IP/port → RDMA device and discover route

Resolve address - `rdma_resolve_addr(cm_id, src, dst, timeout)`

Process CM event (RDMA\_CM\_EVENT\_ADDR\_RESOLVED) -  
`rdma_get_cm_event(echannel, cm_event),`  
`rdma_ack_cm_event(cm_event)`

Resolve route - `rdma_resolve_route(cm_id, timeout)`

Process CM event (RDMA\_CM\_EVENT\_ROUTE\_RESOLVED)

- Server bind address - `rdma_bind_addr(cm_id, server_addr)`
- Server listen - `rdma_listen(cm_server_id, num)`
- Wait and process CM event  
`RDMA_CM_EVNET_CONNECT_REQUEST`  
This gives client id (like TCP conn) from  
`cm_event → id`
- Ack CM event

- Allocate Protection Domain (PD): memory region protections for access control
  - `ibv_alloc_pd(cm_id->verbs)`
- Create Completion Queue (CQ): where completion events (CQEs) are reported after WRs finish. Use a completion channel to get notifications.
  - `ibv_create_comp_channel(cm_id->verbs)`
  - `ibv_create_cq(cm_id->verbs, CQ_CAPACITY, context, comp_channel, 0)`
  - `ibv_req_notify_cq(cq, 0)` //tell CQ that we want notification

- Create Queue Pair (QP): Two queues per endpoint: Send Queue (SQ) and Receive Queue (RQ). The core send/receive structure
  - Set `ibv_qp_init_attr`
  - `rdma_create_qp(cm_id, pd, &qp_init_attr)`


# Connection

- Set parameters – rdma\_conn\_param
- On server side accept - rdma\_accept(cm\_id, &conn\_param)
- On client side connect - rdma\_connect(cm\_id, &conn\_param)

# Memory Region (MR)

# Work Request (WR)

# Scatter-Gather Entry (SGE)



# Memory Registration

- A registered memory buffer the RNIC can read/write
  - Allocate buffer – malloc()/mmap()
  - Register MR: Tell RNIC it can access this memory -  
`ibv_reg_mr(pd, buf, len, flags)`

- SGEs tells the RNIC which buffer in memory the upcoming Work Request (WR) should use.

```
struct ibv_sge sge;
```

```
sge.addr = (uintptr_t)mr->addr; // address of  
buffer
```


```
sge.length = mr->length; // size of buffer
```

```
sge.lkey = mr->lkey; // local access key
```

- Can have multiple SGEs to describe disjoint buffers

```
struct ibv_send_wr send_wr = {0};  
struct ibv_send_wr *bad_wr = NULL;  
send_wr.sg_list = &sge;      // pointer to array of SGEs  
send_wr.num_sge = 1;          // number of SGEs  
send_wr.opcode = IBV_WR_SEND; // or  
IBV_WR_RDMA_WRITE, IBV_WR_RDMA_READ, etc.  
send_wr.send_flags = IBV_SEND_SIGNALED; // ask for CQE  
  
ibv_post_send(qp, &send_wr, &bad_wr);
```

```
struct ibv_recv_wr recv_wr = {0}, *bad_recv =  
NULL;  
  
recv_wr.sg_list = &sge;  
recv_wr.num_sge = 1;  
  
ret = ibv_post_recv(qp, &recv_wr, &bad_recv);
```



C  
Q



RECV  
RECV

C  
Q

- CQE – when you know the WR completes
  - Get notification on CQ channel: there's work completion (WC) in CQ –  
`ibv_get_cq_event(comp_channel,&cq_ptr,  
&context);`
  - Request for more notification -  
`ibv_req_notify_cq(cq_ptr, 0)`
  - Poll WC – `ibv_poll_cq(cq_ptr, wc_num, wc)`
  - Ack CQ events – `ibv_ack_cq_events(cq_ptr, 1)`

# wr\_id

- When having many WRs, which WR is a CQE for?
  - Use wr\_id
  - Set recv\_wr.wr\_id
  - Check wc.wr\_id when receiving CQE later

# Notes

- When SEND, make sure there's at least one posted RECV WR on the other end
- The Receive Queue is FIFO: The first RECV WR you post will be the first one consumed by the next incoming SEND message.
- You can post a RECV again after using it.

# Write/Read remote memory directly

- wr.opcode:  
`IBV_WR_RDMA_WRITE/IBV_WR_RDMA_READ`
- Needs remote raddr + rkey (send by receiver)  
`wr.wr.rdma.remote_addr = remote_base + off;`  
`wr.wr.rdma.rkey = remote_rkey;`

# Disconnect

- Client disconnect - `rdma_disconnect(cm_id)`
- Server wait for disconnect CM event (and ack it) – `RDMA_CM_EVENT_DISCONNECTED`
- Client also wait for disconnect CM event

# Clean up

- Destroy QP – `rdma_destroy_qp(cm_id)`
- Destroy ID – `rdma_destroy_id(cm_id)`
- Destroy CQ – `ibv_destroy_cq(cq)`
- Destroy completion channel – `ibv_destroy_comp_channel(channel)`
- Destroy other event channel –  
`rdma_destroy_event_channel(channel)`
- Destroy PD - `ibv_dealloc_pd(pd)`

# Usefull Resources

- <https://man7.org/linux/man-pages/index.html>

# Lab 3

- Sender (provided framework, your job),  
Receiver (provided)
- Goal: send messages, send file

# Lab 3 Workflow

- Sender setup and initiates connection
- Sender sends file length + hash to receiver
- Receiver sends (addr, len, rkey) of a MR back
- Sender writes the file to that MR, sends a MSG\_DONE when completes
- Receiver compares hash and sends back good/bad
- Sender sends ID
- Receiver sends back secret key